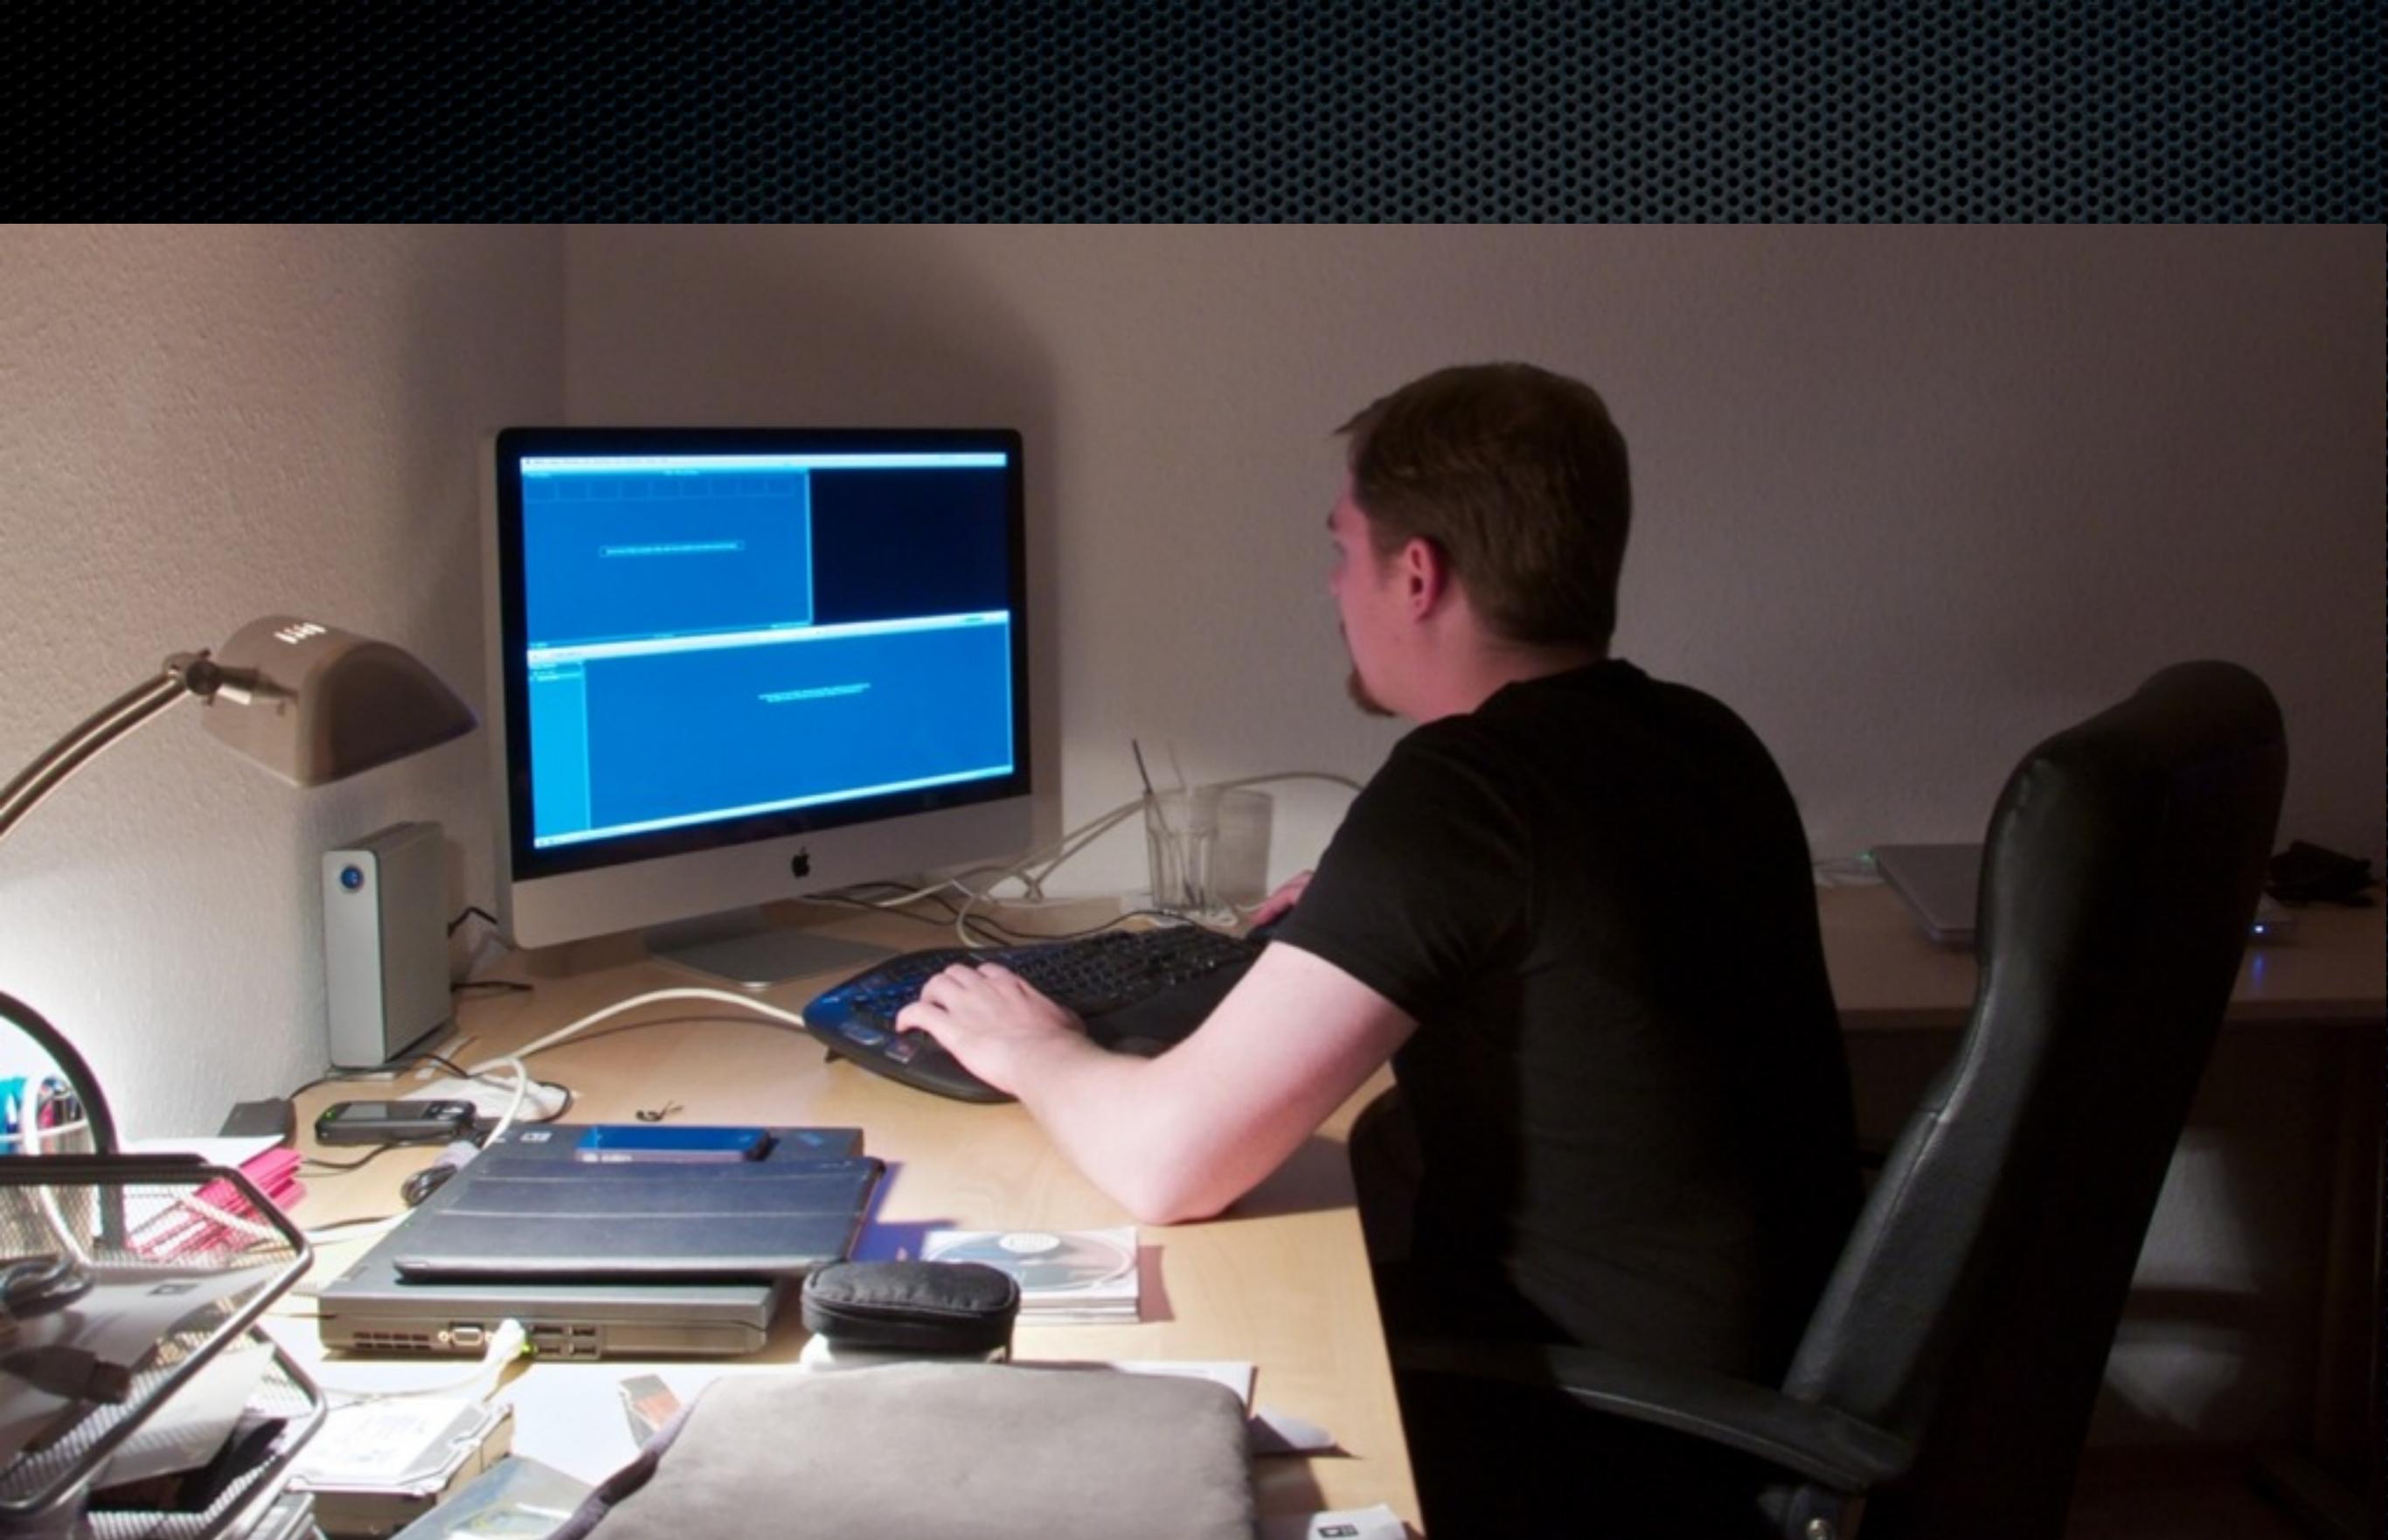
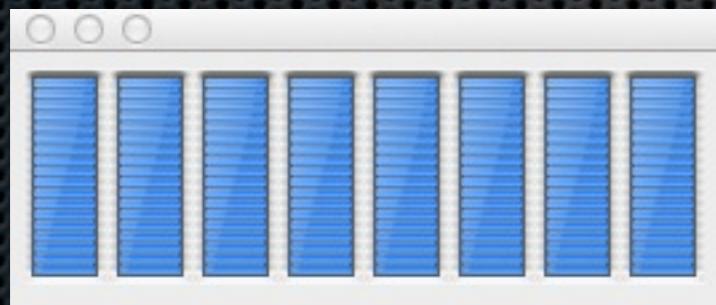


# DTrace & Instruments



<http://www.flickr.com/photos/rob-wei/6156468402/>

# When suddenly...



# Why?

(h)top

Activity Monitor

strace

truss

printf

ps

gdb

vm\_stat

# DTrace!

**“DTrace allows you to ask arbitrary questions about what the system is doing, and get answers.”**

**~ Bryan Cantrill**

**Developed @ Sun  
First Release 2005**

**Solaris, Mac OS X,  
FreeBSD, NetBSD, Oracle  
Linux**

OHAI!

Solaris, Mac OS X,  
FreeBSD, NetBSD, Oracle  
Linux

**(since 10.5)**

# What can it trace?



Trace all the programs in user space

# ...and kernel functions

as well as kernel functions.

# Detour: P\_LNOATTACH

15

On OS X, processes can set this flag to prevent DTrace from attaching to them



Bascially, because of iTunes

# Hello World

```
dtrace -n ':::BEGIN
{
    printf("Hello, world.");
    exit(0);
}'
```

```
dtrace: description ':::BEGIN' matched 1 probe
CPU    ID          FUNCTION:NAME
 0      1          :BEGIN Hello, world.
```

# A more useful example

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

Run with sudo, go make a cup of coffee, and hit Ctrl-C when you come back

# How it works

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

So what does this do?

## enable specified probes

```
dtrace -n syscall::open:entry  
{  
    @num[execname] = count();  
}'
```

`-n` tells dtrace we're specifying our stuff directly on the command line. You can use a file as well, in fact, that's what you'll probably prefer to do.

## probe description

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

This describes the probe we want to match. We don't have to provide the exact name of the probe, we can, in fact, match several probes with one description. More on this later.

```
dtrace -n 'syscall::open:entry
{
    here be actions
    @num[execname] = count();
}'
```

`{` lets DTrace know that between here and `}` are all the actions we want to take when probes fire.

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

**associative array**

`@num` is an associative array called `num`. Note we didn't have to declare it before we used it. Associative arrays can have multiple keys.

There are also scalar variables that simply don't have a prefix.

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

**key for array  
is the executable name**

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

**value is the number of  
times it has been called**

```
dtrace -n 'syscall::open:entry
{
    @num[execname] = count();
}'
```

**‘tis be all the  
action**

`}` lets DTrace know that this is where the fun ends

# Syntax

But that is not all

```
#!/usr/sbin/dtrace -s  
  
/* This is a comment */  
provider : module : function : name  
/ predicate /  
{  
    actions;  
}
```

This is how the general syntax basically looks like

## C-style comments

```
#!/usr/sbin/dtrace -s
```

```
/* This is a comment */
provider:module:function:name
/predicate/
{
    actions;
}
```

C-style comments, no //

```
#!/usr/sbin/dtrace -s
```

```
/* This is a comment */  
provider : module : function : name  
/ predicate /  
{  
    actions;  
}
```

**predicate**  
**filters**

Predicates allow us to specify conditions on when the actions are executed

```
#!/usr/sbin/dtrace -s  
  
/* This is a comment */  
provider : module : function : name  
/ predicate /  
{  
    actions;  
}
```

Also note that you can have more than one probe definition in your file.

```
$macro  
scalar_variable  
@aggregate_array  
self->thread_local  
this->clause_local
```

36

These are the kinds of variables and their scope you have available in DTrace.

self-> are thread-local variables, they are local to the firing probe  
this-> are local to a clause

Variables can be declared as a type, but don't have to be.

```
#!/usr/sbin/dtrace -s  
  
syscall::read:entry  
/execname != "dtrace"/  
{  
    @reads[execname,  
          fds[arg0].fi.pathname] =  
        count();  
}
```

This is a more complicated usage of associative arrays.

This records how often executables access a certain file path. “fds” is an array of structs that contain information about file descriptors, like “fi\_pathname”.

LaunchBar	??/English.lproj/InfoPlist.strings	1
LaunchBar	??/Resources/Exceptions.plist	1
SyncServer	??/Local/cleanup.time	1
Twitter	??/Contents/Info.plist	1
Twitter	??/Resources/Exceptions.plist	1
less	<unknown (not a vnode)>	1
less	??/78/xterm-256color	1
less	??/<unknown (NULL v_name)>/tty	1
mdworker	??/Cookies/Cookies.binarycookies	1
mdworker	??/History/http:%2F%2Ft.co%2FgaGL6M0.webhistory	1
mtmfs	??/Caches/com.apple.FindSystemFiles.plist	1
[...]		
bash	??/libexec/rbenv-version-file	216
git-remote-http	<unknown (not a vnode)>	253
trustevaluation	??/system/mdsDirectory.db	260
Propane	<unknown (not a vnode)>	273
bash	??/libexec/rbenv-version-name	288
git	??/.git/HEAD	330
nVALT	<unknown (not a vnode)>	460
FindSystemFiles	??/Contents/Info.plist	492
WebProcess	<unknown (not a vnode)>	787
sh	<unknown (not a vnode)>	951
bash	<unknown (not a vnode)>	2704

pid	PID of firing process
ppid	PPID of firing process
errno	Last errno encountered by firing process
execname	Name of process firing the probe
arg0 ... argX	64bit integers value returned from probe
args[0]...args[X]	Typed values returned from a probe

\$pid	PID of dtrace process
\$ppid	PPID of dtrace process
\$cwd	Current working directory
\$0...\$9	Arguments to DTrace (\$0 is script name)
\$target	This will resolve to a PID of a process using the -p or -c switch
[\$e](g u)id	effective and real group or user id

Macros are expanded by DTrace when it processes the input file. You can explicitly convert a macro to a string by adding an additional “\$” in front, e.g. \$\$pid

count()	Returns the number of times called
avg()	computes arithmetic average
sum()	adds the input to the aggregation
max()/min()	Returns max/min of the inputs provided
quantize()	generate a power of 2 distribution of the input
printf()	It's printf as we know and love/hate it

Selection of functions available.

“printa” is also very useful to print out an aggregation.

```
copyin(pointer, size)
copyinstr(pointer_to_str)
```

Important: When probes pass you addresses to pointers, they are relative to that process, not dtrace. `copyin` copies whatever the pointers point to, and make it available in your dtrace process. `copyinstr` is specifically for NULL-terminated strings.

# DTrace Probes

# Many, many probes

If you have DTrace, then you'll have a lot of probes to choose from

```
dtrace -l
```

Pipe this to `less` or a file, it's a lot of output.

# ~300.000 probes available

46

Right now, on my system, there are about 300k probes available for me to, well, probe.

dtrace	Pseudo-probes like BEGIN, END, ...
syscall	Entry and return points for every syscall
profile	Time-based probes that fire on ticks, seconds or fractions
io	Disk-I/O related probes
proc	process creation, forking and signaling
objc	Information about ObjC messages
ruby	Ruby memory/GC information
pid<pid>	Allows you to trace userland functions

Selection of probe providers.

syscall entry args are the same as the syscall gets. syscall return arg0 are the return value of the syscall.

[http://wikis.sun.com/  
display/DTrace/Providers](http://wikis.sun.com/display/DTrace/Providers)

# Embedding your own

```
provider primes {  
    /* Start of the prime calculation */  
    probe primecalc_start(long prime);  
  
    /* End of the prime calculation */  
    probe primecalc_done(long prime, int  
        isprime);  
  
    /* Exposes the size of the table of existing  
    primes */  
    probe primecalc_tablesize(long tableszie);  
};
```

Simple example: We have a program to calculate primes, and we want three DTrace probes.

First, we write a prime\_probes.d file like this.

```
dtrace -o prime_probes.h -h -s prime_probes.d
```

```
#include <stdio.h>

long primes[1000000] = { 3 };
long primecount = 1;

int main(int argc, char **argv)
{
    long divisor = 0; long currentprime = 5; long isprime = 1;

    while (currentprime < 1000000)
    {
        isprime = 1;

        for(divisor=0;divisor<primecount;divisor++)
        {
            if (currentprime % primes[divisor] == 0) { isprime = 0; }

        if (isprime)
        {
            primes[primecount++] = currentprime;
            printf("%ld is a prime\n",currentprime);
        }
        currentprime = currentprime + 2;
    }
}
```

```

#include <stdio.h>
#include "prime_probes.h"

long primes[1000000] = { 3 };
long primecount = 1;

int main(int argc, char **argv)
{
    long divisor = 0; long currentprime = 5; long isprime = 1;

    while (currentprime < 1000000)
    {
        isprime = 1;
        PRIMES_PRIMECALC_START(currentprime);
        for(divisor=0;divisor<primecount;divisor++)
        {
            if (currentprime % primes[divisor] == 0) { isprime = 0; }
        }
        PRIMES_PRIMECALC_DONE(currentprime, isprime);
        if (isprime)
        {
            primes[primecount++] = currentprime;
            PRIMES_PRIMECALC_TABLESIZE(primecount);
            printf("%ld is a prime\n",currentprime);
        }
        currentprime = currentprime + 2;
    }
}

```

```
clang -c primes.c  
clang -o primes primes.o
```

```
#!/usr/sbin/dtrace -s

#pragma D option quiet

primes*:::primecalc-start { self->start = timestamp; }

primes*:::primecalc-done
/arg1 == 1/
{
    @times["prime"] = sum(timestamp - self->start);
}

primes*:::primecalc-done
/arg1 == 0/
{
    @times["nonprime"] = sum(timestamp - self->start);
}

END
{
    normalize(@times,1000000);
    printa(@times);
}
```

```
» ./primes
```

(in another shell)

```
» sudo dtrace -s timing.d  
^C
```

prime	3344
nonprime	385675795

```
if (PRIMES_PRIMECALC_START_ENABLED()) {  
    PRIMES_PRIMECALC_START(currentprime);  
}
```

If you want to be a good citizen, you use this. It also allows you to calculate additional information to pass to the probe.

# Instruments

# Demo

# apropos dtrace

See all the scripts that use DTrace to provide you with information

# DTrace

DYNAMIC TRACING IN ORACLE® SOLARIS,  
MAC OS X, AND FREEBSD



Brendan Gregg • Jim Mauro  
Foreword by Bryan Cantrill